

Video grader: X3D-M with shared 3- and 5-class heads (2.99 M parameters)



Each res stage stacks X3D blocks (1×1×1 expand → 3×3×3 depthwise → squeeze-excite → 1×1×1 project, residual); time stays at 16 frames.
Loss: class-weighted CE on each head, summed. Fix: the Kinetics head's Softmax is removed, so both losses see logits.